

Implications of using LLMs to explain bereavement to children

Introduction

Communicating the death of a loved one to a child is one of the most difficult conversations a caregiver may ever face. The challenge is not only emotional but also cognitive: How do we explain the permanence of death to a mind still grasping basic concepts of time, emotion, and abstract thinking? As large language models (LLMs) like ChatGPT become increasingly embedded in everyday digital communication, some users have begun turning to these tools in moments of emotional vulnerability—including moments of grief. This paper explores whether LLMs might be used to support or simulate the task of introducing the concept of death to a child, and what the psychological, ethical, and communicative implications of that use might be.

Using a mix of qualitative prompt-response analysis and a review of developmental and communication literature, this paper argues that while LLMs can generate emotionally sensitive, age-appropriate responses when prompted carefully, their use in this context should be approached with caution. While these tools may offer linguistic clarity and emotional scaffolding for caregivers under stress, they cannot—and should not—replace the presence, judgment, and empathy of a human caregiver when dealing with childhood grief.

Literature Review

Understanding Death in Childhood Communication

Children make sense of death differently at different stages of their development. According to Piaget's theory, younger children (under 7) often think death is reversible or temporary because they're still learning how abstract concepts work (Piaget, 1952; Speece & Brent, 1996). As they grow—typically between ages 7 to 11—they start to understand that death is permanent, universal, and inevitable. Still, these understandings don't develop in a vacuum. Age is important, but so are personal experiences, family beliefs, and cultural background.

Nagy (1948) offers a helpful breakdown: first, young children may think of death like a long trip or sleep; then, they may think it's something that happens to other people but can be avoided; and finally, by around age nine or ten, they start to really grasp that death is final and affects everyone. Meanwhile, Bluebond-Langner's (1978) work with terminally ill children showed that even young kids can develop deep, complex understandings of death—especially when the people around them talk about it openly or give off nonverbal clues. Her research reminds us that kids often pick up more than adults think, and their understanding is shaped just as much by their environment as by their cognitive stage.

The Role of Language and Emotion in Grief Communication

When it comes to grief communication, language is everything. It's not just about giving information—it's about helping children feel safe, heard, and less alone. That's why many experts, like Schonfeld and Demaria (2016), recommend using clear, simple, and direct language. Euphemisms like “went to sleep” can be more harmful than helpful, especially for kids who might then become scared of going to sleep themselves.

Christ (2000) points out that adults play two roles when they're breaking hard news to children: they explain what happened, and they model how to manage the emotions that come with it. Silverman (2000) and Worden (1996) both emphasize that grieving isn't just one conversation—it's an ongoing process. Children will keep asking questions, and their understanding will grow over time. Creating a space where those conversations can continue is essential.

It's also important to be mindful of the kinds of language used, especially when religion or spirituality comes into play. Balk and Corr (2009) found that abstract or religious language can either soothe or confuse, depending on what the child already believes or has been taught. For instance, saying “Grandma is watching over us from heaven” may be comforting to some, but to others, it could be confusing or upsetting. That kind of nuance is hard to replicate with an LLM unless it's guided by someone who really knows the child.

Large Language Models in Emotionally Sensitive Contexts

LLMs like ChatGPT are increasingly being used in emotionally complex situations, including mental health, education, and caregiving. While these models can generate text that sounds empathic, they aren't actually feeling anything—they're just predicting what words should come next based on patterns in their training data. As Gnewuch et al. (2022) explain, this means they can simulate empathy, but they don't really understand the emotion behind the words. Still, many users have said they feel supported or comforted by talking to them, even knowing it's not a real person.

To better measure how well LLMs can handle emotional scenarios, Chen et al. (2024) developed EmotionQueen, a benchmark that tests how models perform on tasks like recognizing key emotional events or mixed emotions. Their findings show that LLMs can handle straightforward emotional content fairly well, but they struggle when things get more complicated—like when a child expresses guilt or conflicting emotions. This is a major limitation when it comes to grief, where responses need to be nuanced and highly sensitive.

What's also concerning is that many AI tools now being marketed for children—like educational bots or storytelling assistants—haven't been evaluated for use in grief or trauma contexts. And there aren't clear rules or ethical standards guiding how these tools should respond when a child brings up something serious like death. Without that guidance, there's a real risk that a child could get a confusing or even harmful response.

In short, the existing research shows that while LLMs can help caregivers generate clear, age-appropriate explanations for children, they aren't capable of navigating the full emotional depth and relational complexity that these conversations require. Their responses should be seen as a tool—not a replacement—for the kind of care, connection, and trust that only people can provide.

Methodology and Example Prompts

This paper uses qualitative content analysis of LLM-generated responses to simulate the communicative scenario. Prompts were entered into OpenAI's ChatGPT (GPT-4) with different parameters: age of the child, relationship to the deceased, and emotional state of the caregiver.

Prompt 1: Parent telling a 6-year-old that Grandma has died

Prompt:

"Can you help me explain to my 6-year-old that their grandmother has died? I want to be honest but gentle. I don't want to scare them, but I want them to understand she's not coming back."

LLM Output (excerpt):

"You can say something like: 'I have something very sad to tell you. Grandma was very sick, and her body stopped working. That means she died. When someone dies, it means we can't see them or talk to them anymore, but we can still remember all the love they gave us. It's okay to feel sad or confused. I'm here with you, and we can talk about it anytime you want.'"

Analysis:

This response is emotionally sensitive, avoids euphemisms, and uses age-appropriate metaphors ("her body stopped working"). It anticipates emotional confusion and offers reassurance. However, the model does not ask clarifying questions about the child's beliefs or previous experiences with death prior to giving a response

Analysis and Implications

Affordances of LLMs

LLMs can offer emotionally supportive phrasing, context-specific suggestions, and rapid access to grief communication strategies. For overwhelmed or grieving caregivers, they may serve as a source of language scaffolding or emotional guidance. The model's capacity to adapt tone and rephrase in response to follow-up queries allows for personalization. LLMs also provide a low-risk environment for emotional rehearsal. Caregivers can try different ways of explaining death without the pressure of an audience. This allows them to work through their own emotional blocks, fears, or uncertainties. For example, writing out a message with the model's help might clarify a caregiver's internal confusion about what they want to say, and why.

Risks and Limitations

However, LLMs also carry significant limitations. They do not assess the emotional state of the child or the caregiver. Their responses may be inconsistent or culturally inappropriate. Moreover, they may inadvertently perpetuate misinformation or inappropriate metaphors (e.g., "sleeping" as a euphemism for death), potentially complicating a child's understanding of mortality.

There are significant ethical concerns about outsourcing emotionally charged responsibilities—like explaining the death of a loved one to a child—to machines. First, such delegation risks creating emotional distance between caregivers and children. These conversations, though difficult, offer opportunities for connection, empathy, and modeling healthy emotional expression. When a caregiver hands this moment over to an AI, even partially, they may unintentionally signal avoidance or discomfort, rather than offering presence and support. Over time, this could undermine the child's trust in their caregiver as a safe source of comfort and information.

Additionally, LLMs lack situational awareness and contextual nuance, making them potentially risky in complex family dynamics, religious contexts, or in cases involving traumatic or sudden death. If relied upon too heavily, they could produce scripted responses that seem out of place or emotionally tone-deaf, especially if caregivers do not know how to prompt them carefully. There is also the danger of accidental harm: a poorly phrased AI suggestion could evoke fear, guilt, or confusion in a child who is already vulnerable. Unlike trained grief counselors, LLMs cannot detect when a message has caused distress or redirect the conversation based on nonverbal cues or emotional escalation.

Finally, dependence on LLMs for emotionally significant communication could contribute to a broader societal shift in how we relate to each other. If digital tools become common intermediaries in moments of grief, there is a risk of diminishing our collective capacity to engage in difficult conversations, weakening the interpersonal bonds that are so vital in moments of mourning. Relying on LLMs may allow adults to bypass the emotional labor involved in grieving with a child. While it can be helpful to receive support in structuring thoughts or finding words, the emotional work of sitting with a child's pain and answering follow-up questions is essential. Ethical concerns emerge when technology is used to short-circuit that process. Grief, especially for young children, is not resolved through single explanations, but through repeated, sensitive, and responsive conversations—something no AI can sustainably provide.

Developmental Sensitivity

LLM-generated content must be rigorously vetted for developmental alignment. A five-year-old cannot process metaphysical reassurance or abstract concepts like legacy or peace. Yet some LLM responses introduce such terms without adequate simplification. While models can be prompted to adjust tone or vocabulary, the onus is on the user to guide the output—a risk if the caregiver is emotionally dysregulated.

Conclusion

While large language models can offer support in drafting developmentally appropriate, empathetic explanations of death for children, they cannot replace the emotional presence, follow up support, intuition, and responsibility of human caregivers. These models may be best viewed as tools to support—not supplant—grief conversations. To responsibly integrate LLMs in such contexts, future development must include more rigorous safeguards, clearer disclaimers, and potentially hybrid models that pair LLM output with professional-reviewed guidance. Ultimately, the question is not just whether LLMs *can* explain death to children, but whether we *should* rely on them to do so—and what might be lost when we do.

Final Bibliography

Balk, D. E., & Corr, C. A. (2009). *Children's Encounters with Death, Bereavement, and Coping*. Springer.

Bluebond-Langner, M. (1978). *The Private Worlds of Dying Children*. Princeton University Press.

Chen, Y., Yan, S., Liu, S., Li, Y., & Xiao, Y. (2024). EmotionQueen: A Benchmark for Evaluating Empathy of Large Language Models. *Findings of the Association for Computational Linguistics: ACL 2024*, 2149–2176. <https://aclanthology.org/2024.findings-acl.128/>

Christ, G. H. (2000). *Healing Children's Grief: Surviving a Parent's Death from Cancer*. Oxford University Press.

Gnewuch, U., Morana, S., & Maedche, A. (2022). Designing empathic conversational agents. *Journal of the Association for Information Systems*, 23(3), 608–636.

Nagy, M. H. (1948). The child's theories concerning death. *Journal of Genetic Psychology*, 73(1), 3–27.

OpenAI. (2023). *GPT-4 System Card*. <https://openai.com/research/gpt-4-system-card>

Schonfeld, D. J., & Demaria, T. (2016). Supporting the grieving child. *Pediatrics*, 138(3), e20162147.

Speece, M. W., & Brent, S. B. (1996). The development of children's understanding of death. *Death Studies*, 20(1), 29–52.

Worden, J. W. (1996). *Children and Grief: When a Parent Dies*. Guilford Press.

.